


Concurrent Computing

Lecturers: Prof. Majid Mirmehdi (majid@cs.bris.ac.uk)
Dr. Tilo Burghardt (tilo@cs.bris.ac.uk)
Dr. Daniel Page (page@cs.bris.ac.uk)

Web: <http://www.cs.bris.ac.uk/Teaching/Resources/COMS20001>



LECTURE 2

TOWARDS CONCURRENT PROGRAMMING IN πC

Recap: The natural world is not serial ... ☺

- ...NATURE is massively concurrent !
 - natural networks tend to be continuously evolving, yet they are robust, efficient and long-lived
 - Concurrency is one of nature's core design mechanisms – and one of ours!





- in many cases **computing models phenomena of the real world**
 - computers are built as part of the physical world and can harvest natural concurrency for their own performance
 - concurrency can often help simplifying the modelling of systems


Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 2

Recap: Multiprocessors & Multi-Core Revolution

- **Multiprocessors**
(collection of communicating processors)
 - speed advantage by physically parallelised computation





- **Multi-Memory Systems**
 - local, CPU-associated memory essential regardless of programming model
 - however, connectivity model affects specific performance tradeoffs



Connectivity is Critical: Bus or Point-to-Point?

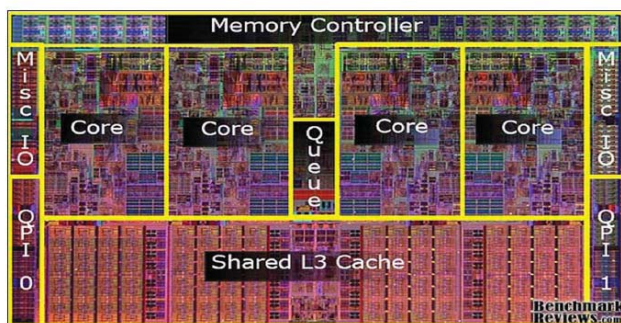
- traditional design: front-side-bus (FSB)
 - each processor has to compete for access
 - multitude of processors/resources result in bottleneck
- ways forward: **localised memory** and **on-chip networks (switch)**
 - multiple simultaneous point-to-point (P2P) connections between cores & resources (much like end-to-end 'channels')



Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 3

Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 4


Example 1: Intel i7 Architecture



- each processor has its own dedicated memory using integrated memory controller (IMC)
- P2P QuickPath onchip net for cross-core access + snoop traffic

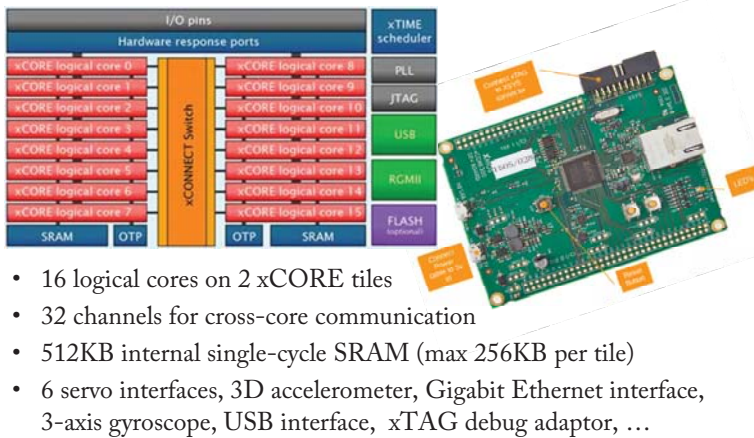
Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 5

Example 2: Future IBM OnChip Optical Interconnects



Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 6


Example 3: XMOS xCore200 Explorer Kit



Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 7

Why learn XC?

- built around multi-threading and point-to-point **channel communication** model (...thus, in line with current hardware design trends...)
- compiles directly to drive **multi-core hardware**
- familiar C syntax**, yet semantic similarities to classic parallel languages such as Occam
- theoretically grounded in **process algebra CSP**, which can be used to reason about (usually basic) XC programs

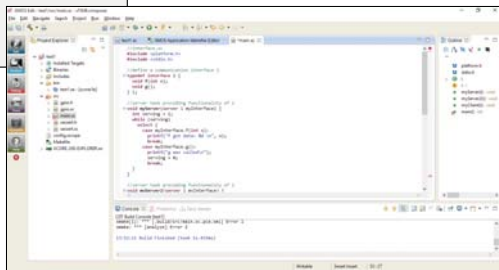


Getting Started: `hello.xc` & xTimeComposer IDE

```
// hello.xc
#include <stdio.h>

int main(void) {
    printf("Hello!\n");
    return 0;
}
```

- installed in Linux lab MVB2.11
- freely available for download for Linux, Windows and MacOS (www.xmos.com/support/tools)




Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 9

So far in C...

Sequential Processes: Deterministic Control Flow

- given an input, a single sequential process (=thread) produces a **single sequence of memory state changes** deriving the output
- XC: every **basic block {}** is treated as a sequential process with a strict order of execution (...first do this, next this ...)




XC example: sequential process its memory trace for input (2,4)

Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 10

Where C cannot go – a wish list

- EXPLICIT PARALLELISM**
we want to execute several statements **in parallel** **on different cores** to gain a speed advantage over sequential execution (...trading temporal spread for spatial spread of computation)
- EXPLICIT COMMUNICATION**
we want to **channel messages between cores/threads** to synchronise several concurrent computations
- EXPLICIT CONTROL**
we want to **control the physical location** of execution and storage to minimise data transfers and effectively use local resources (...compactness under programmer control)

XC Concurrent Execution – PAR statement




Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 11

Burghardt & Mirmehdi | 2014 | Concurrent Computing COMS20001 | Computer Science | University of Bristol | Lecture 2 / 12

Non-Deterministic Control Flow

- given an input, a set of concurrent processes (=threads) produces **one out of many possible sequences** of memory state changes deriving an output ('implicit choice' during runtime)

→ XC: every statement/sub-block in a **par{}** block is treated as an independent process



Program Example: Execution in Parallel

```
// par.xc
#include <platform.h>
#include <stdio.h>

void hello(int threadNo);

// main starting two tasks in parallel
int main(void) {
    par {
        hello(0); //start first thread in parallel
        hello(1); //start second thread in parallel
    } // wait until both threads have terminated
    return 0;
}

// function to print message
void hello(int threadNo){
    printf("Hello from thread %d.\n", threadNo);
}
```

Key Concept: Process Construction

Two ways of combining processes into a compound process:

- **SEQUENTIAL concatenation** in a block **{}**
 - returns when its last component process finishes
 - executed one after the other (...implicit in XC...)
- **PARALLEL composition** in a block **par{}**
 - written order of components is irrelevant
 - returns when all its component processes have returned

Any process, compound or just a single statement, ...

- **starts** (i.e. thread is instantiated),
- **performs** a number of actions (i.e. thread runs)
- and then **may finish/terminate** (i.e. thread returns to caller)

Revisited: Concurrency vs. Parallelism

CONCURRENCY...

...is concerned with **non-deterministic composition** of processes (i.e. program components)

PARALLELISM...

...is concerned with **exploiting independencies** among the sub-computations of a deterministic computation

NO COMPILER... is available today that automatically turns a sequential process into a set of communicating, concurrent processes that optimally exploit independencies among the sub-computations

→ programmers need to understand concurrent programming paradigm to exploit & support emerging physical parallelism

Example: Execution on Different Physical Tiles

```
// helloworld.xc
#include <platform.h>
#include <stdio.h>

void hello(int tileNo);

// main starting two tasks in parallel on different tiles
int main(void) {
    par {
        on tile[0] : hello(0); //start on tile 0
        on tile[1] : hello(1); //start on tile 1
    }
    return 0;
}

// function to print message
void hello(int tileNo){
    printf("Hello from tile %d.\n", tileNo);
}
```

Example: Combining XC with C Sources & Timers

```
//partxc.xc
#include <platform.h>
#include <stdio.h>

extern void hello(int tileNo);

int main(void) {
    par {
        on tile[1] : hello(1);
        on tile[0] : hello(0);
    }
    return 0;
}

//delays execution
void delay(uint delay)
{
    uint time, tmp;
    //define a timer
    timer t;
    //read current state of timer
    t->> time;
    //trigger when timer has moved on the delay no of ticks
    when timerafter ( time + delay ) :> tmp;
}
```

```
//partc.c
#include <stdio.h>
#include <platform.h>

extern void delay(uint delay);

void hello(int tileNo){
    delay((3-tileNo)*1000);
    printf("Hello from tile %d.\n",tileNo);
}
```

Example: Interfaces for Single Client-Server Setups

```
//interface.xc
#include <platform.h>
#include <stdio.h>

//define a communication interface i
typedef interface i {
    void f(int x);
    void g();
} i;

//server task providing functionality of i
void myServer(serve i myInterface) {
    int serving = 1;
    while (serving)
        select {
            case myInterface.f(int x):
                printf("f got data: %d \n", x);
                break;
            case myInterface.g():
                printf("g was called\n");
                serving = 0;
                break;
        }
} ...

...
//client task calling function
//of task 2
void myClient(client i myInterface) {
    myInterface.f(2);
    myInterface.f(1);
    myInterface.g();
}

//main starting two threads
//calling over an interface
int main() {
    interface i myInterface;
    par {
        myServer(myInterface); //only1server
        myClient(myInterface); //only1client
    }
    return 0;
}
```



```
Console Problems Task Viewer
<terminated> test1.xe [xCORE Application] xrun
f got data: 2
f got data: 1
g was called
```

Outlook to Lecture 3



Message Passing / Channel Communication

*or "How a thread running on Core#1
can talk to a thread running on Core#2?"*